

HybridPerfopticon: Query Visualization for Hybrid Distributed Databases

Brandon Haynes Shrainik Jain
University of Washington
{bhaynes, shrainik}@cs.washington.edu

ABSTRACT

Hybrid databases have recently gained a lot of attention because of the promise they show in several prominent use cases in scientific data management. A hybrid distributed database system usually consists of multiple underlying databases with different data models. In order to better understand such systems we present a query profiling and visualization tool for queries authored over these hybrid systems. Our approach is to extend an existing visualization tool for distributed databases by adding support for multiple underlying database systems.

Author Keywords

Hybrid Databases, Query Visualization

ACM Classification Keywords

H.5.2 Graphical user interfaces (GUI): Miscellaneous

INTRODUCTION

Hybrid databases are an emerging area in database research. A hybrid database comprises of dissimilar systems fused together in various ways. Stonebraker *et al.* [9] recently observed that the “one size fits all” era is over, and that the demands of big data management warrant a more federated approach to database design. Lim *et al.* in Cyclops [7] offered a top-down approach that combined compatible systems into a monolithic whole. Polybase [5], on the other hand, took a bottom-up approach where disparate systems were more loosely connected. To aid understanding, development and debugging of these new systems, we need better visualization tools as well.

Many profiling systems exist that target queries in a single database system. For example, SQL Server, Postgres, MySQL all have a basic “explain” functionality, which displays a query plan in a human-readable (but often unfriendly) format. Most ‘explain’-type tools lack deeper analysis features such as the ability to view performance over the duration of the query. Further compounding this problem is the limited availability of tools available in a distributed database

context, where operator-worker assignments, data skew, and similar problems are much more important.

Some tools, however, have been recently introduced that attempt to fill this need in a distributed database environment. Gprof [2], is a basic visualization tool for system development; however, it tends to be oriented toward developers. Dapper [1] is Google’s production distributed systems tracing tool designed to aid in distributed system development, but it lacks support for distributed query visualization. Finally, Ambrose [4] is a framework designed to visualize map-reduce jobs. While it offers many of the features that one would like to see in a robust distributed database visualization tool it currently targets only the map-reduce space.

All of these tools lack a complete visualization story (*i.e.* where the tool is useful across both users and system developers) which has resulted in a scarcity of debugging and visualization in the hybrid distributed database space. PerfOpticon [8] seems to be the most advanced system that addresses the profiling requirements for a distributed database system and is targeted to both audiences.

This paper presents an extension to PerfOpticon to support these emerging use cases, where users require robust profiling of plans that co-occur across multiple database systems. These present a number of unique challenges, such as how to display non-relational features (e.g., array chunks) in a framework that has heretofore targeted only tuple-based relational systems. Additionally, a number of corollary challenges arise including coordinating the display of multiple sub-plans, visualizing data transfer between systems, and gathering profiling data from multiple sources.

For the purpose of demonstration, we chose a hybrid distributed database with SciDB [3] and Myria [6] as its component systems. SciDB’s data-model is completely different from Myria, making it critical to generate semantics for communication between these systems. HybridPerfOpticon, our extension to PerfOpticon, highlights this communication by helping visualize the overall hybrid query plan, tuple flow among operators, skew in distribution of workloads among system and internal query time statistics while being completely compatible with the existing design.

In the next section, we present the architecture of the system, which details the steps necessary to orchestrate hybrid plan execution, creation, and serving logs (for visualization purposes). This is followed by some results from the

use case analysis so far. We conclude by looking at some future extensions to HybridPerfoticon.

HYBRIDPERFOPTICON ARCHITECTURE

Our approach to extending Perfoticon to support hybrid queries involved three major additions:

- Generating HybridPerfoticon-compatible SciDB plans
- Creating HybridPerfoticon by extending Perfoticon to recognize components from varying underlying DBMSs
- Creating hybrid plans and exposing them to HybridPerfoticon

In this section we describe each of these changes in detail and highlight the trade-offs in our design decisions.

Generating SciDB Plans & Metrics

In the Myria DBMS, each operator is instrumented to store profiling data. When Perfoticon needs this profiling data, a HTTP request is generated, a query is performed over these data, and the aggregated results are returned.

While we could take a similar route in SciDB by instrumenting each operator and responding to aggregation requests, this would be a time-consuming task and would decrease the performance of SciDB itself. Instead, we elected an alternative approach: since the relevant profiling metadata is already present in the SciDB logs, we generated our performance metrics by scraping these log files. While this resulted in decreased performance on startup (approximately ten seconds per 100MB of logging data), our approach has the added benefit of requiring no changes to the SciDB source code. We expect that this approach would allow for a low-barrier introduction of similar systems, assuming that the salient metrics exist in the log files.

The architecture of our log-scraping process is illustrated in Figure 2. The entry point to this process is the `SciDBParser` class, which is responsible (via the `SciDBLog` class) for converting raw log entries into sequences of statements. Statements associated with a given query are aggregated into a physical plan, and the operators for that plan are identified and hierarchically assembled within that plan. During the log interpretation process we also identify performance metrics associated with each executed query and associate those results with the plan.

An illustration of the log-scraping process is listed in Figure 1. In this log fragment (abbreviated for readability), the `SciDBLogParser` identifies the beginning of the query execution process, identifies (and interprets) the associated physical plan, associates a scatter/gather operation with an underlying operator, and notes the time at which the query is committed. Each of these entries are converted into an equivalent statement, and the useful metadata for each statement is embedded therein.

Extending Perfoticon to Support Additional DBMSs

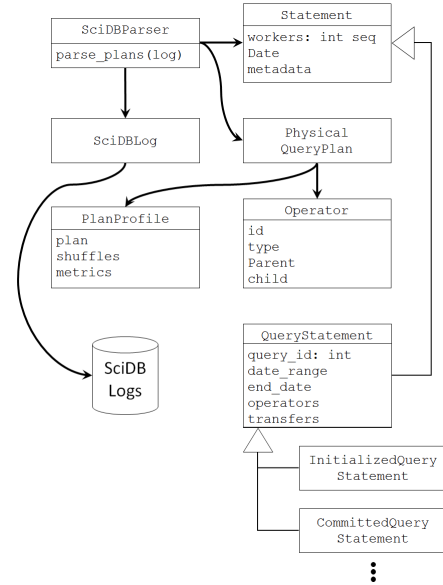


Figure 2. Architecture of the SciDB log parsing components.

In order to extend Perfoticon to support queries across multiple DBMSs, we first explored the use of existing extension points. Since Perfoticon is not designed to be directly extendable (e.g., via a plugin-style architecture), extensive modification to the framework would be required to support the kinds of hybrid plans we intended to visualize. Instead, we elected an alternative approach where a substantial portion of the necessary logic was hidden behind the HTTP server (a “hybrid HTTP server”, described in detail below) to which Perfoticon connects for plan and profiling metadata.

While this approach greatly reduced the modifications required to realize our desired functionality, some changes were nonetheless required to Perfoticon itself. For example, it is necessary to be able to identify the system of origin for each plan fragment; this is important not only for visualization purposes (e.g., to color fragments by system type) but also functionally. For example, the hybrid HTTP server must know from which system a fragment originated in order to retrieve profiling metrics.

A number of similar minor modifications were required throughout the Perfoticon framework. By way of example, SciDB fragments required a SciDB-specific query identifier be attached to each fragment; we envision that additional metadata would be required in order to support other systems.

Hybrid HTTP server

Given that HybridPerfoticon continues to expect HTTP responses in the same form as Perfoticon, our hybrid HTTP server needed to do a substantial amount of manipulation to compose hybrid plans without violating this requirement. A high-level architectural overview of our approach is illustrated in Figure 3.

For each incoming request, our server first distinguished between three cases:




Figure 4. Detailed profiling results for the SciDB fragment in Figure 6.

Algorithm 1 Combining Myria and SciDB plans

```

function CREATE-HYBRID-PLAN(id, urlMYRIA)
1: plan = HTTP-GET(id, urlMYRIA)
2: for each fragmentMYRIA ∈ plan do
3:   fragmentMYRIA.origin = "Myria"
4:   for each opMYRIA ∈ fragmentMYRIA do
5:     if opMYRIA.type is SciDBScan then
6:       planSciDB = load(opMYRIA.parent.id)
7:       for each fragmentSciDB ∈ planSciDB do
8:         fragmentSciDB.origin = "SciDB"
9:         plan.fragments += fragmentSciDB
10:      for each opSciDB ∈ fragmentSciDB do
11:        if opSciDB.type is MyriaSave then
12:          opSciDB.child.id = opMYRIA.id
13: return plan

```

queries lags a bit due to many API calls to the hybrid plan server. We are exploring possible ways to fix this.

CONCLUSIONS & FUTURE WORK

In this work we described HybridPerfopticon, a visualization tool for hybrid queries that cross DBMS boundaries. A tool that does so is critically needed, given that the performance profiles of this class of queries are not currently well-understood. Accordingly, we expect that this tool will serve as a useful means by which the hybrid DBMS space may be explored.

While our approach of using existing log files to generate performance metrics is certainly not novel, the use of log file

```

SciDB {
  save(
    redimension(
      load('auto-mpg.csv', -2, 'csv'),
      <mpg: double, type:string>[id=0:*, 1, 0]),
      'intermediate', -1, 'csv+');
}

Myria {
  mpgs = load("intermediate",
    csv(schema(id:int,
      mpg:float,
      type:string)));
  prices = load("auto-prices.csv",
    csv(schema(type:string,
      price:int)));
  dollars_per_mpg = [from mpgs, prices
    where mpgs.type = prices.type
    emit id, type, price / mpg];
  store(dollars_per_mpg, result);
}

```

Figure 5. A hybrid query used as input to HybridPerfopticon. See Figure 6 for generated hybrid plan.

parsing to quickly add a second DBMS to the visualization framework represents an opportunity to extend the reach of HybridPerfopticon to a wide-ranging set of systems. Future work should evaluate the applicability of this approach, and if generalizable, additional attention should be paid to identifying the optimal interface by which logging metadata may be exposed to HybridPerfopticon.

However, while our log-parsing approach allowed us to quickly expose profiling data for a second DBMS inside HybridPerfopticon, more work remains to be done. In particular, the hybrid HTTP server needs a far more robust caching sub-

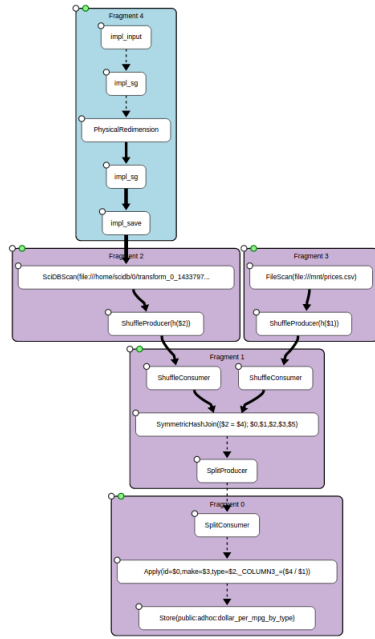


Figure 6. A hybrid plan as visualized by HybridPerfopticon. This hybrid query begins with a SciDB scan/redimension, after which the resulting array is transferred to Myria for joining on a second dataset.

system; in particular, increasingly large logging files will quickly lead to an intractable startup time. This would be easily remedied by ensuring that the results of the parsing operation were themselves logged to a DBMS for quick subsequent retrieval.

Finally, future work should evaluate how to best degrade in the face of differing levels of profiling granularity. For example, we produced SciDB profiling metrics at the millisecond level, whereas Myria profiles at nanosecond granularity. A hybrid visualization framework should be able to gracefully reconcile both; similarly, we expect that some systems will not have support for some profiling areas and that the framework should evolve to handle these cases.

REFERENCES

1. Google dapper. <http://research.google.com/pubs/pub36356.html>.
2. Gprof. <https://sourceware.org/binutils/docs/gprof/>.
3. Scidb by paradigm4. <http://www.paradigm4.com/>.
4. Twitter ambrose. <https://github.com/twitter/ambrose>.
5. D. J. DeWitt, A. Halverson, R. Nehme, S. Shankar, J. Aguilar-Saborit, A. Avanes, M. Flasz, and J. Gramling. Split query processing in polybase. In *Proceedings of the 2013 ACM SIGMOD International Conference on Management of Data*, pages 1255–1266. ACM, 2013.

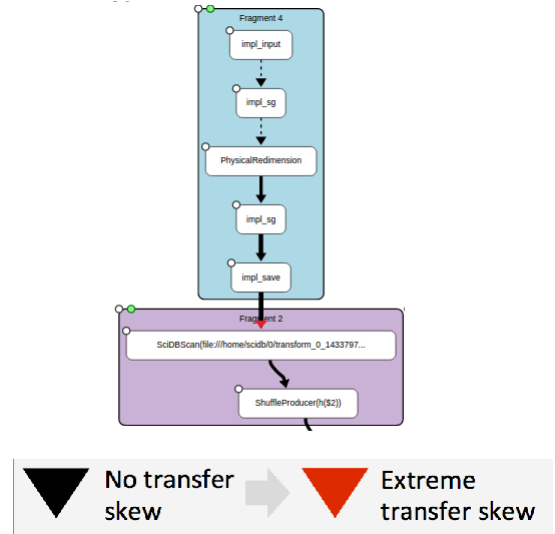


Figure 7. Visualizing skew in inter-database transfers

6. D. Halperin, V. Teixeira de Almeida, L. L. Choo, S. Chu, P. Koutris, D. Moritz, J. Ortiz, V. Ruamviboonsuk, J. Wang, A. Whitaker, et al. Demonstration of the myria big data management service. In *Proceedings of the 2014 ACM SIGMOD international conference on Management of data*, pages 881–884. ACM, 2014.
7. H. Lim, Y. Han, and S. Babu. How to fit when no one size fits. In *CIDR*, volume 4, page 35, 2013.
8. D. Moritz, D. Halperin, B. Howe, and J. Heer. Perfopticon: Visual query analysis for distributed databases. In *Computer Graphics Forum (EuroVis), Cagliari, Italy*, volume 34, 2015.
9. M. Stonebraker, S. Madden, D. J. Abadi, S. Harizopoulos, N. Hachem, and P. Helland. The end of an architectural era:(it’s time for a complete rewrite). In *Proceedings of the 33rd international conference on Very large data bases*, pages 1150–1160. VLDB Endowment, 2007.